

# Linear Regression with Gradient Descent

## 1 Introduction

Linear regression is a fundamental machine learning technique used to model the relationship between a dependent variable  $y$  and an independent variable  $x$  as a linear equation  $y = wx + b$ , where  $w$  is the weight (slope) and  $b$  is the bias (intercept). This report describes the implementation of linear regression using gradient descent on a small dataset of 4 points, loaded from a CSV file. The implementation, written in Python, minimizes the mean squared error (MSE) loss function to find the optimal  $w$  and  $b$ . We analyze the results and the effect of the learning rate on the algorithm's performance.

## 2 Dataset

The dataset consists of 4 points, loaded from `lr.csv`. Based on the plot, the points are approximately:

$$(30, 60), \quad (40, 80), \quad (60, 120), \quad (80, 160).$$

These points suggest a linear relationship  $y \approx 2x$ , with some offset.

## 3 Algorithm Description

Gradient descent is an iterative optimization algorithm that minimizes the MSE loss:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2,$$

where  $y_i$  is the true value,  $\hat{y}_i = wx_i + b$  is the predicted value, and  $n$  is the number of data points.

The algorithm updates  $w$  and  $b$  using the gradients of the MSE with respect to each parameter:

$$\begin{aligned} \frac{\partial \text{MSE}}{\partial w} &= -\frac{2}{n} \sum_{i=1}^n x_i (y_i - \hat{y}_i), \\ \frac{\partial \text{MSE}}{\partial b} &= -\frac{2}{n} \sum_{i=1}^n (y_i - \hat{y}_i). \end{aligned}$$

The update rules are:

$$w \leftarrow w - \eta \cdot \frac{\partial \text{MSE}}{\partial w}, \quad b \leftarrow b - \eta \cdot \frac{\partial \text{MSE}}{\partial b},$$

where  $\eta$  is the learning rate.

A stopping threshold is used to halt the algorithm if the change in cost between iterations is below a threshold (e.g.,  $10^{-6}$ ).

## 4 Implementation

The Python implementation includes the following steps:

1. **Data Loading:** Read  $x$  and  $y$  values from `lr.csv`, expecting 4 points.
2. **MSE Loss:** Compute the MSE using:

$$\text{MSE} = \frac{1}{n} \sum (y_{\text{true}} - y_{\text{pred}})^2.$$

### 3. Gradient Descent:

- Initialize  $w = 0.1$ ,  $b = 0.01$ .
- For up to 1000 iterations:
  - Compute predictions:  $\hat{y}_i = wx_i + b$ .
  - Compute MSE.
  - If the cost change is below  $10^{-6}$ , stop.
  - Compute gradients using the above formulas.
  - Check for numerical instability (gradients or parameters too large).
  - Update  $w$  and  $b$ .
  - Store  $w$ ,  $b$ , and cost for plotting.
- Return final  $w$ ,  $b$ , and histories.

### 4. Visualization: Plot the data points with the fitted line and the cost vs. weights.

The parameters used are:

- Learning rate:  $\eta = 0.0001$ .
- Maximum iterations: 1000.
- Stopping threshold:  $10^{-6}$ .

## 5 Results

The gradient descent algorithm converged to:

$$w = 2.04, \quad b = 1.84.$$

The resulting linear model is:

$$y = 2.04x + 1.84.$$

### 5.1 Regression Plot

The left plot (Figure 1) shows the 4 data points and the fitted line. The line closely matches the data, indicating a good fit. The points lie near the line, confirming the linear relationship  $y \approx 2x$ , with a small positive intercept.

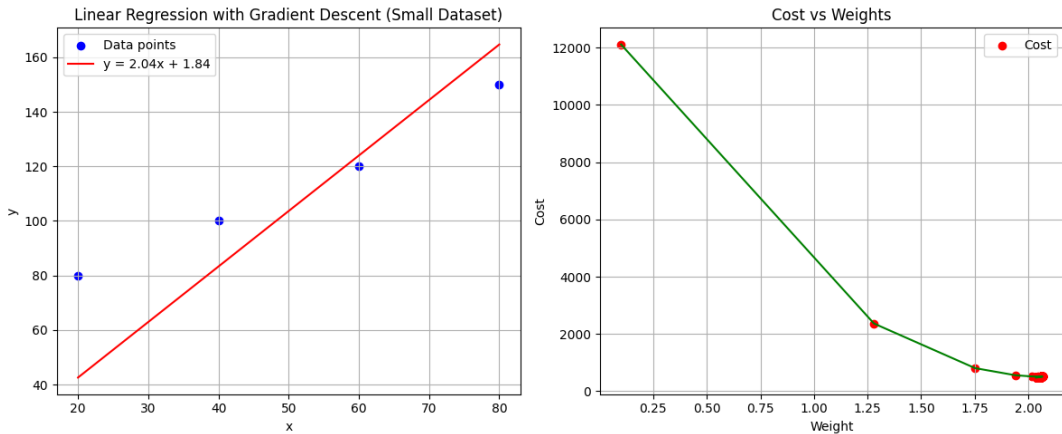


Figure 1: Linear regression plot (left) and cost vs. weights (right).

### 5.2 Cost vs. Weights

The right plot (Figure 1) shows the MSE cost as a function of the weight  $w$ . The cost decreases from approximately 12000 to below 2000, with  $w$  increasing from 0.1 to 2.04. The trajectory is smooth, indicating stable convergence, though the small learning rate results in many iterations to reach the minimum.